Differentially Private Distributed Matrix Multiplication: Fundamental Accuracy-Privacy Trade-Off Limits

Ateet Devulapalli[†], Viveck R. Cadambe[†], Flavio P. Calmon[‡], Haewon Jeong[‡]

[†]Pennsylvania State University, [‡] Harvard University

Abstract—The classic BGW algorithm of Ben Or, Goldwasser and Wigderson for secure multiparty computing demonstrates that secure distributed matrix multiplication over finite fields is possible over $2t + 1$ computation nodes, while keeping the input matrices private from every t colluding computation nodes. In this paper, we develop and study a novel coding formulation to explore the trade-offs between computation accuracy and privacy in secure multiparty computing for real-valued data, even with fewer than $2t + 1$ nodes, through a differential privacy perspective. For the case of $t = 1$, we develop achievable schemes and converse arguments that bound ϵ — the differential privacy parameter that measures the privacy loss — for a given accuracy level. Our achievable coding schemes are specializations of Shamir secret sharing applied to real-valued data, coupled with appropriate choice of evaluation points. We develop converse arguments that apply for general additive noise based schemes.

Index Terms—Differential privacy, privacy-utility tradeoff, mean square error, secure multiparty computation, coded computing, distributed matrix multiplication.

I. INTRODUCTION

The task of accurate and efficient distributed data processing while preserving data privacy is among the most important engineering problems in modern machine learning. The desire to keep data private inevitably requires the source adding some noise to the data before sharing it with the computation nodes. Secure multiparty computing (MPC) is a paradigm that ensures that data remains private from any t computing nodes[1], yet it guarantees accurate computation of functions of the data [1]. The celebrated BGW algorithm [2] provides a method to perform *information-theoretically* private computation of a wide class of functions building on Shamir's secret sharing technique [3], which, in turn, builds on Reed Solomon codes.

Consider two random matrices $\mathbf{A}, \mathbf{B} \in \mathbb{F}^{L \times L}$, where $\mathbb{F}$ is a field, and a set of P computation nodes. Let $\mathbf{R}_i, \mathbf{S}_i, i = 1, 2, \dots, t$ be statistically independent $L \times L$ random matrices. In Shamir's secret sharing, node i receives inputs $\tilde{\mathbf{A}}_i = p_1(x_i), \tilde{\mathbf{B}}_i = p_2(x_i)$, where, $x_1, x_2, \dots, x_P$ are distinct non-zero scalars and $p_1(x), p_2(x)$ are matrix-valued polynomials:

$$p_1(x) = \mathbf{A} + \sum_{j=1}^{t} \mathbf{R}_j x^j, p_2(x) = \mathbf{B} + \sum_{i=1}^{t} \mathbf{S}_j x^j.$$

[1]In this paper, we study the setting called honest-but-curious or semi-honest adversary [1] in Secure MPC terminology.

If the field $\mathbb{F}$ is finite, and the entries of $\mathbf{R}_i, \mathbf{S}_i, i = 1, 2, \dots, t$ are chosen randomly i.i.d. uniformly over the elements of the field, then the input to any subset $\mathcal{S}$ of t nodes is independent of the data $(\mathbf{A}, \mathbf{B})$. If node i computes $\tilde{\mathbf{A}}_i + \tilde{\mathbf{B}}_i$, then observe that the sum $\mathbf{A} + \mathbf{B}$ — which is the constant in the polynomial $p_1(x) + p_2(x)$ — can be recovered from the computation output of any $t + 1$ of the P nodes by polynomial interpolation. Observe, similarly, that $\tilde{\mathbf{A}}_i \tilde{\mathbf{B}}_i$ can be interpreted as an evaluation at $x = x_i$ of the degree $2t$ polynomial $p_1(x)p_2(x)$, whose constant term is $\mathbf{AB}$. Thus, the matrix product $\mathbf{AB}$ can be recovered from any $2t + 1$ nodes via polynomial interpolation. The BGW algorithm uses the above coding scheme to perform secure MPC for the universal class of computations that can be expressed as sums and products, while maintaining (perfect) data privacy among every set of t nodes. Because of its universality, the BGW algorithm forms the basis of several secure MPC protocols. Notably, an overhead of $2t + 1$ computation nodes (i.e., $t + 1$ redundant nodes) are required to keep the data private from any t computation nodes and perform multiplications, as compared to mere data access as in Shamir secret sharing, or addition/aggregation wherein the computation output can be recovered from $t + 1$ computation nodes. In fact, for more complex functions, the overhead can be prohibitively large inevitably leading to multiple communication rounds [2], or more redundant nodes [4]. For instance, for polynomial computations (of which multiplication is a special case), the number of redundant nodes required to enable single round computations can scale linearly as the polynomial degree [4].

We study the canonical and fundamental operation of multiplication, and aim to answer the following question: Can codes be developed to enable a set of fewer than $2t + 1$ nodes to compute the matrix product while keeping the data (matrices) private from any t nodes? While impossibility results preclude this possibility if we aim for the dual goals of exact recovery of the matrix product *and* perfect privacy, we study the question through a novel coding formulation that allows for approximations on both fronts, and thereby enables a study of their trade-off.

For machine learning applications that inevitably operate over real-valued data, approximate computation of the output typi-

cally suffices. Further, a prevalent[2] paradigm for data privacy in machine learning applications in practice is *differential privacy (DP)* [7], which aims to keep small perturbations of the data private. In particular, it requires that the extent of *privacy loss* (see Sec. II) be bounded by a non-negative parameter ϵ; the smaller the ϵ, the lesser the privacy loss. The case of matrices $\mathbf{A}, \mathbf{B}$ being independent of the nodes' input corresponds to the special case of $\epsilon = 0^3$. While the DP framework allows us to tune the degree of privacy, to effectively use the framework in practice, it is important to understand how to set parameter ϵ based on the application at hand (See [8]); our paper aims to bring this understanding to the context of secure matrix multiplication.

A. Summary of Contributions

Our main contribution is an explicit analytical characterization of the trade-off between computation accuracy and privacy for the case of $t = 1$, that is, where the data is kept private from every single computation node in the system. We consider the following problem. Assume that a computation node gets an input of the form $\mathbf{A} + \mathbf{R}, \mathbf{B} + \mathbf{S}$ and multiplies them, where $\mathbf{R}, \mathbf{S}$ are random noise matrices that independent of $(\mathbf{A}, \mathbf{B})$ designed to ensure data privacy. The goal of the decoder is to recover an estimate $\tilde{\mathbf{C}}$ of the product $\mathbf{AB}$ from N computation outputs at a certain accuracy level. The noise $\mathbf{R}, \mathbf{S}$ should ensure that the data $(\mathbf{A}, \mathbf{B})$ is ϵ-differentially private (ϵ-DP) from the input to every $t = 1$ computation node.

For $N < 3$, we characterize via an achievable coding scheme and a converse, the trade-off between mean square error $||\tilde{\mathbf{C}} - \mathbf{AB}||_F$ and the DP parameter ϵ. For both the achievable coding scheme and the converse, we follow a two step procedure. We first develop bounds on the mean square error in terms of the second moments of Frobenius norm and singular values of the noise matrices $\mathbf{R}, \mathbf{S}$. In a second step, we bound the DP parameter ϵ in terms of these second moments, applying a specific distribution for our achievable scheme, and bounding ϵ over all distributions for the converse. For the case where $\mathbf{A}, \mathbf{B}$ are scalars, our first step translates to a *tight* characterization between the mean square error and the standard deviations of $\mathbf{R}, \mathbf{S}$. Our achievable scheme is a specialization of the Shamir secret sharing technique applied for real numbers with a careful choice of evaluation points, followed by a DP analysis. Our converse makes only mild assumptions on the structure of codes, applies to a general class of schemes for additive noise.

B. Related Work

Differential Privacy and Secure MPC: Several prior works are motivated like us to reduce computation and communication overheads of secure MPC by connecting it with the less

stringent privacy guarantee offered by DP. References [9]–[12] provide methods to reduce communication overheads for sample aggregation algorithms, label private training private record linkage, private distributed median computation. In comparison we aim to reduce the overhead of t redundant nodes for multiplication and use the DP framework to develop an analytical accuracy-privacy tradeoff.

Coded Computing: The emerging area of coded computing enables the study of codes for secure computing that enable data privacy. Our framework resonates with the coded computing approach, as we abstract the algorithmic/protocol related aspects into a master node, and highlight the role of the error correcting code in our model. Coded computing has been applied to study code design for secure multiparty computing in [4], [13]–[18]. These references effectively extend the standard BGW setup by imposing memory constraints on the nodes, or other constraints, that effectively disable each node storing information equivalent to the entire data sets. Under the imposed constraints, these references develop novel codes and characterize regimes for exact computation and perfect privacy. In particular, codes for secure MPC over real-valued fields have been studied in [13], [19] extending the ideas of [4] to understand the loss of accuracy due to finite precision. In particular, reference [13] casts the effect of finite precision in a privacy-accuracy tradeoff framework. In contrast to all previous works in coded computing geared towards secure MPC, we operate below the threshold of perfect recovery, and characterize the trade-off between them. Our incorporation of differential privacy for this characterization is a novel aspect of our set up. We do not impose any memory constraints on the nodes, and imposition of such constraints can lead to interesting areas of future study.

Privacy-Utility Trade-offs. There is a fundamental trade-off between DP and utility (see [20]–[22] for examples in machine learning and statistics). The optimal ϵ-DP noise-adding mechanism for a target moment constraint on the additive noise was characterized in [23]. For approximate DP, near-optimal additive noise mechanisms under ℓ_1-norm and variance constraints were recently given in [24]. Our converse makes use of a lower-bound on the variance of an ϵ-DP noise mechanism that is looser, albeit simpler than [23, Thm. 7].

II. SYSTEM MODEL AND PROBLEM STATEMENT

Notations: We define $[n] \triangleq \{1, 2, \cdots, n\}$. We use bold fonts for vectors and matrices. We define $(\mathbf{x})_i$ to be the i^{th} component of a vector $\mathbf{x}$ and $(\mathbf{X})_{k,l}$ be the $(k, l)^{\text{th}}$ element of a matrix $\mathbf{X}$. Denote $\kappa(\mathbf{X}), ||\mathbf{X}||_2$ and $||\mathbf{X}||_F$ to be the minimum singular value, ℓ_2 norm and Frobenius norm of a matrix $\mathbf{X}$ respectively. We use $X \sim \mathbb{Q}$ to say that the random variable X has the probability distribution $\mathbb{Q}$.

A. System Model

We consider a computation system with P computation nodes. $\mathbf{A}, \mathbf{B} \in \mathbb{R}^{L \times L}$ are random matrices, and node $i \in [P]$ receives:

$$\tilde{\mathbf{A}}_i = \mathbf{A} + \mathbf{R}_i, \quad \tilde{\mathbf{B}}_i = \mathbf{B} + \mathbf{S}_i$$

---

[2]See, for example, Google's Tensorflow Privacy Framework [5], [6].

[3]The Shamir secret sharing in fact can be applied for real-valued data as well. Specifically, by allowing evaluation points x_i to grow arbitrarily large, the DP parameter ϵ can be made arbitrarily small and still allow for perfectly accurate decoding of the matrix product from any $2t + 1$ nodes.

where $\mathbf{R}_i, \mathbf{S}_i \in \mathbb{R}^{L \times L}$ are random matrices such that $(\mathbf{R}_1, \mathbf{R}_2, \ldots, \mathbf{R}_P, \mathbf{S}_1, \mathbf{S}_2, \ldots, \mathbf{S}_P)$ is statistically independent of $(\mathbf{A}, \mathbf{B})$. We denote by $\mathbf{R}, \mathbf{S} \in \mathbb{R}^{L \times PL}$, the following:

$$\mathbf{R} = \begin{bmatrix} \mathbf{R}_1 & \mathbf{R}_2 & \ldots & \mathbf{R}_P \end{bmatrix}$$
$$\mathbf{S} = \begin{bmatrix} \mathbf{S}_1 & \mathbf{S}_2 & \ldots & \mathbf{S}_P \end{bmatrix}.$$

In this paper we assume no shared randomness between $\mathbf{R}, \mathbf{S}$, i.e., they are statistically independent: $\mathbb{P}_{\mathbf{R},\mathbf{S}} = \mathbb{P}_{\mathbf{R}} \mathbb{P}_{\mathbf{S}}$. We denote by $\mathcal{P}_{\mathbf{R},\mathbf{S}}$ as the set of all possible joint distributions of $\mathbf{R}, \mathbf{S}$ where $\mathbf{R}, \mathbf{S}$ are independent.

For $i \in [P]$, computation node i outputs

$$\tilde{\mathbf{C}}_i = \tilde{\mathbf{A}}_i \tilde{\mathbf{B}}_i. \tag{2.1}$$

A decoder receives the computation output of an arbitrary set $\mathcal{S}$ of N nodes and performs a map: $d_{\mathcal{S}} : (\mathbb{R}^{L \times L})^{|\mathcal{S}|} \to \mathbb{R}^{L \times L}$ that is linear over $\mathbb{R}$. That is, the decoder outputs:

$$\widetilde{\mathbf{C}}_{\mathcal{S}} = d_{\mathcal{S}}(\tilde{\mathbf{C}}_i|_{i \in \mathcal{S}}) = \sum_{i \in [|\mathcal{S}|]} w_{i,\mathcal{S}} \tilde{\mathbf{C}}_i \tag{2.2}$$

where the coefficients $w_{i,\mathcal{S}} \in \mathbb{R}, i \in [|\mathcal{S}|]$ specify the linear map $d_{\mathcal{S}}$. A (P, N) *coding scheme* for positive integers $P \geq N$ consists of the joint distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}} \in \mathcal{P}_{\mathbf{R},\mathbf{S}}$, and the decoding maps $\prod_{\mathcal{S} \subseteq P : |\mathcal{S}| = N} \{d_{\mathcal{S}} : (\mathbb{R}^{L \times L})^{|\mathcal{S}|} \to \mathbb{R}^{L \times L}\}$. The performance of a coding scheme is measured by two metrics: privacy and accuracy.

Remark 1. The standard secure multiparty computing set up assumes that $P = N$. We keep our system model general and allow P to be larger than N. When P is larger than N, the developed schemes have the benefit of tolerance to $P - N$ failures/stragglers, in addition to data privacy and accurate computations.

Privacy of a (P, N) coding scheme

Definition 2.1. (*t-node ϵ-DP*) Let $\epsilon \geq 0$. The distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ satisfies t-node ϵ-DP if, for arbitrary matrices $\mathbf{A}_0, \mathbf{B}_0, \mathbf{A}_1, \mathbf{B}_1 \in \mathbb{R}^{L \times L}$ that satisfy $\left\| \begin{bmatrix} \mathbf{A}_0 \\ \mathbf{B}_0 \end{bmatrix} - \begin{bmatrix} \mathbf{A}_1 \\ \mathbf{B}_1 \end{bmatrix} \right\|_{\max} \leq 1$,

$$\frac{\mathbb{P}\left(\mathbf{Y}_{\mathcal{T}}^{(0)} \in \mathcal{A} \right)}{\mathbb{P}\left(\mathbf{Y}_{\mathcal{T}}^{(1)} \in \mathcal{A} \right)} \leq e^{\epsilon}, \tag{2.3}$$

for all subsets $\mathcal{T} \subseteq [P], |\mathcal{T}| = t$, for all subsets $\mathcal{A} \subset \mathbb{R}^{2L \times tL}$ in the Borel σ-field, where

$$\mathbf{Y}_{\mathcal{T}}^{(\ell)} \triangleq \begin{bmatrix} \mathbf{A}_\ell + \mathbf{R}_{i_1} & \mathbf{A}_\ell + \mathbf{R}_{i_2} & \ldots & \mathbf{A}_\ell + \mathbf{R}_{i_{|\mathcal{T}|}} \\ \mathbf{B}_\ell + \mathbf{S}_{i_1} & \mathbf{B}_\ell + \mathbf{S}_{i_2} & \ldots & \mathbf{B}_\ell + \mathbf{S}_{i_{|\mathcal{T}|}} \end{bmatrix}, \ell = 0, 1$$

where $\mathcal{T} = \{i_1, i_2, \ldots, i_{|\mathcal{T}|}\}$.

We denote by $\mathcal{P}_{\mathbf{R},\mathbf{S}}^{\epsilon,t}$, the set of all possible joint distributions $\mathbb{P}_{\mathbf{R},\mathbf{S}} \in \mathcal{P}_{\mathbf{R},\mathbf{S}}$ that satisfy t-node ϵ-DP. Note that (2.3) depends only on the joint distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ and does not depend on the distributions of $\mathbf{A}, \mathbf{B}$, since the definition applies for arbitrary vectors $\mathbf{A}_0, \mathbf{B}_0, \mathbf{A}_1, \mathbf{B}_1$ — that is, those that are not necessarily drawn from $\mathbb{P}_{\mathbf{A},\mathbf{B}}$.

Accuracy of a (P, N) coding scheme

The main goal of this paper is to characterize the trade-off between privacy and accuracy of estimation of the matrix-product $\mathbf{AB}$. In particular, we develop schemes that guarantee a certain level of DP (i.e., a certain value of parameter ϵ), irrespective of the distribution of the inputs. It is, however, necessary (and standard, see [13], [17], [25], [26]) to account for the data distribution and its parameters when evaluating the accuracy of coding schemes. The accuracy guarantees of the coding schemes developed in this paper rely on the following key assumptions:

Assumption 2.1. $\mathbf{A}$ and $\mathbf{B}$ are statistically independent random matrices. Moreover, there is a parameter $\eta > 0$ such that:

$$\mathbb{E}\left[||\mathbf{A}||_F^2\right] = \mathbb{E}\left[||\mathbf{B}||_F^2\right] \leq \eta.$$

We measure the accuracy of a coding scheme via the mean square error. Specifically, we define:

Definition 2.2 (Linear Mean Square Error (LMSE)). For a (P, N) coding scheme Γ consisting of joint distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ decoding maps $\prod_{\mathcal{S} \subseteq P : |\mathcal{S}| = N} \{d_{\mathcal{S}} : (\mathbb{R}^{L \times L})^{|\mathcal{S}|} \to \mathbb{R}^{L \times L}\}$, the LMSE for subset $\mathcal{S} \subseteq [P], |\mathcal{S}| = N$ is defined as:

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) = \mathbb{E}[||\mathbf{AB} - \widehat{\mathbf{C}}_{\mathcal{S}}||_F^2]. \tag{2.4}$$

where $\widehat{\mathbf{C}}_{\mathcal{S}}$ is defined in (2.2). The LMSE of the coding scheme Γ is defined to be:

$$\mathsf{LMSE}(\Gamma) = \max_{\mathcal{S}} \mathsf{LMSE}_{\mathcal{S}}(\Gamma).$$

It is worth noting that the expectation in the above definition is over the joint distributions of the random variables $\mathbf{A}, \mathbf{B}, \mathbf{R}, \mathbf{S}$. In particular, the accuracy of a coding scheme can depend on the parameters[4] of the joint distribution of $\mathbf{A}, \mathbf{B}$. Sometimes we will explicitly denote the distribution in the LMSE notation as: $\mathsf{LMSE}_{\mathcal{S}}^{\mathbb{P}_{\mathbf{A},\mathbf{B}}}(\Gamma)$ or $\mathsf{LMSE}^{\mathbb{P}_{\mathbf{A},\mathbf{B}}}(\Gamma)$.

Definition 2.3 (Optimal Mean Square Error (LMSE$^*(P, N, \epsilon, t)$)**.**

$$\mathsf{LMSE}^*(P, N, \epsilon, t) = \inf_{\Gamma} \mathsf{LMSE}(\Gamma) \tag{2.5}$$

where the infimum is over the set of all possible (P, N) coding schemes Γ whose joint distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ satisfies t-node ϵ-DP.

The goal of this paper is to characterize $\mathsf{LMSE}^*(P, N, \epsilon, t)$. It is a simple exercise to verify that, if for $N \geq 2t + 1$, coding schemes used by the BGW algorithm achieve $\mathsf{LMSE}^*(P, N, \epsilon, t) = 0$ for all $\epsilon > 0$. That is, perfect privacy[5] and perfect accuracy are achievable for $N \geq 2t + 1$ nodes. Thus, we aim to characterize $\mathsf{LMSE}^*(P, N, \epsilon, t)$ for $N \leq 2t$.

---

[4]It can be readily verified from the LMSE definition that the accuracy simply depends on the means, variances and pairwise correlations of all the random variables involved in $\mathbf{A}, \mathbf{B}, \mathbf{R}_i|_{i=1}^P, \mathbf{S}_i|_{i=1}^P$.

[5]More precisely, ϵ can be made arbitrarily small by adding Laplacian noise of correspondingly large variances, and yet the LMSE can be kept 0.

III. SUMMARY OF RESULTS

The main contribution of this paper is the characterization of an explicit tradeoff between accuracy (LMSE) and privacy (ϵ) for distributed matrix multiplication for the case[6]where $t = 1, N = 2$. The key to our approach is to utilize the variance of the noise as proxy metric for DP, and develop a sharp relation between privacy and accuracy under this metric. Then, the obtained results are translated to bounds on the privacy-accuracy trade-off for ϵ-DP.

We present two technical results. The first is an achievability result that shows that there exists a (P, N) coding scheme with random variables $(\mathbf{R}, \mathbf{S})$ with $\mathbb{E}[||\mathbf{R}_i||_F^2], \mathbb{E}[||\mathbf{S}_i||_F^2] \geq \sigma_{\mathsf{Ach}}^2, \forall i \in [P]$, such that

$$\mathsf{LMSE}(\Gamma) \leq \frac{\eta^2}{\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2}\right)^2} + \Delta$$

for every $\Delta > 0$. The second is a converse that states that, for *any* $\mathcal{S} \subseteq [P], |\mathcal{S}| = 2$:

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) \geq \frac{\eta^2}{\left(1 + \frac{\eta}{\sigma_{\mathsf{Con}}^2}\right)^2}.$$

so long as the minimum singular values of both $\mathbf{R}_i, \mathbf{S}_i$ are lower bounded by σ_{Con}^2 in expectation, that is: $\mathbb{E}[\kappa(\mathbf{R}_i)^2], \mathbb{E}[\kappa(\mathbf{S}_i)^2] \geq \sigma_{\mathsf{Con}}^2, \forall i \in [P]$.

The parameters σ_{Ach}^2 and σ_{Con}^2 intuitively determine the (minimum) variance of the noise added to the inputs at the nodes, and therefore they indirectly control the degree of privacy. In particular, larger values of σ_{Ach} and σ_{Con} corresponds to greater amount of privacy, and correspondingly poorer LMSE. We next discuss details behind the achievability and converse stated above. Our discussions also include bounds implied on the LMSE in terms of the DP parameter ϵ. We provide some proof sketches here; all missing theorem proofs and details can be found in Appendix A.

A. Achievability

Theorem 3.1. *Let* $\mathbf{\Lambda}, \mathbf{\Theta}$ *be* $L \times L$ *independent zero-mean random matrices with i.i.d. entries each with a variance of* $1/L^2$. *For any* $\Delta, \sigma_{\mathsf{Ach}} > 0$ *there exist scalars* $u_i, i \in [P]$ *with* $|u_i| \geq \sigma_{\mathsf{Ach}}$ *such that, if* $\mathbf{R}_i = u_i\mathbf{\Lambda}, \mathbf{S}_i = u_i\mathbf{\Theta}, i \in [P]$, *then there is a* $(P, N = 2)$ *coding scheme* Γ *with distribution* $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ *such that, for every* $\mathbb{P}_{\mathbf{A},\mathbf{B}}$ *satisfying Assumption 2.1,*

$$\mathsf{LMSE}(\Gamma) \leq \frac{\eta^2}{\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2}\right)^2} + \Delta. \tag{3.1}$$

Proof Sketch (Missing details in Appendix A). For simplicity of notation, we sketch the argument for the subset $\mathcal{S} =$

$\{1, 2\}$; the same argument readily extends to an arbitrary two-element subset of $[P]$. As stated in theorem statement, we show an achievable scheme for the subset of distributions $(\mathbf{R}, \mathbf{S})$,

$$\mathbf{R} = \begin{bmatrix} u_1 & \dots & u_P \end{bmatrix} \otimes \mathbf{\Lambda}, \mathbf{S} = \begin{bmatrix} u_1 & \dots & u_P \end{bmatrix} \otimes \mathbf{\Theta}.$$

Node $i \in [P]$ gets evaluations as,

$$\tilde{\mathbf{A}}_i = \mathbf{A} + \mathbf{\Lambda}u_i, \tilde{\mathbf{B}}_i = \mathbf{B} + \mathbf{\Theta}u_i,$$

where $|u_i| \geq \sigma_{\mathsf{Ach}}, \forall i \in [P]$. For convenience of illustration we drop the dependence on $\mathcal{S}$ for decoding weights, i.e., $w_{1,\mathcal{S}}, w_{2,\mathcal{S}}$ will be written as w_1, w_2 respectively. Thus,

$$\widehat{\mathbf{C}}_{\mathcal{S}} = w_1\tilde{\mathbf{A}}_1\tilde{\mathbf{B}}_1 + w_2\tilde{\mathbf{A}}_2\tilde{\mathbf{B}}_2.$$

Then, it can be shown that from the independence of $\mathbf{A}, \mathbf{B}, \mathbf{\Lambda}, \mathbf{\Theta}$ and from the theorem hypothesis that $\mathbb{E}[\mathbf{\Lambda}], \mathbb{E}[\mathbf{\Theta}] = \mathbf{0}, E[||\mathbf{\Lambda}||_F^2], E[||\mathbf{\Theta}||_F^2] = 1$ that:

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) = \mathbb{E}[||\mathbf{AB} - \widehat{\mathbf{C}}_{\mathcal{S}}||_F^2]$$

$$= (w_1 + w_2 - 1)^2\eta^2 + 2\eta(w_1u_1 + w_2u_2)^2 + (w_1u_1^2 + w_2u_2^2)^2.$$

Then minimizing the above expression over w_1, w_2 and then substituting back, we get the following expression,

$$\min_{\substack{w_1, w_2 \\ u_1 \neq u_2}} \mathsf{LMSE}_{\mathcal{S}}(\Gamma) = \frac{2\eta^2u_1^2u_2^2}{2u_1^2u_2^2 + \eta(u_1 + u_2)^2 + 2\eta^2} \tag{3.2}$$

so long as u_1, u_2 are distinct. By choosing *distinct* $u_i, i \in [P]$ arbitrarily close to σ_{Ach}, we obtain, for any $\Delta > 0$,

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) \leq \frac{\eta^2}{\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2}\right)^2} + \Delta$$

$\square$

It is worth noting that the achievable coding scheme is indeed Shamir secret sharing over real field with an appropriate choice of evaluation points. An intriguing aspect of the theorem proof is that the choice of evaluation points u_i is arbitrarily close to σ_{Ach}. Consider the case where $u_2 = \sigma_{\mathsf{Ach}}$, and examine the LMSE with respect to u_1. When u_1 is *equal* to σ_{Ach}, the computation of both nodes $1, 2$ are identical, and this would lead to a poor mean square error. But even a small deviation translates to a near optimal choice of u_1. Consequently, the minimum LMSE is, in fact, a discontinuous function of u_1 at $u_1 = \sigma_{\mathsf{Ach}}$.

We translate the result of Theorem 3.1 to ϵ-DP by restricting $\mathbf{\Lambda}, \mathbf{\Theta}$ to independent Laplace distributions. We implement the widely used Laplace noise distribution here, as it gives a simple achievable scheme that is readily extended to matrices. Our results can potentially be improved, esp. for $L = 1, 2$ by applying the optimal noise distribution under variance constraints studied in [23].

Theorem 3.2. *Let* $\mathbf{\Lambda}, \mathbf{\Theta}$ *be independent zero-mean random matrices with i.i.d. Laplacian distributed entries each with a*

---

[6]The case of $t = 1, N = 1$ is simple to analyze along the arguments of this paper, and is presented in Appendix B .

variance of $1/L^2$. For any $\Delta > 0, \epsilon \geq 0$ there exist scalars $u_i, i \in [P]$ such that, if $\mathbf{R}_i = u_i \mathbf{\Lambda}, \mathbf{S}_i = u_i \mathbf{\Theta}, i \in [P]$, then there exists a $(P, N = 2)$ coding scheme with distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}} \in \mathcal{P}_{\mathbf{R},\mathbf{S}}^{\epsilon,1}$ such that, for every $\mathbb{P}_{\mathbf{A},\mathbf{B}}$ satisfying Assumption 2.1:

$$\mathsf{LMSE}(\Gamma) \leq \frac{\eta^2}{\left(1 + \frac{\eta \epsilon^2}{8L^6}\right)^2} + \Delta.$$

The proof given in Appendix A uses the standard argument that Laplacian mechanisms satisfy ϵ-DP [7].

B. Converse

We derive converse results that lower bound the LMSE for a fixed level of privacy. Similar to our approach to achievability, we first derive a lower bound in Theorem 3.3 in terms of the expected singular values of the noise distributions.

Theorem 3.3. *For any $(P, N = 2)$ code Γ whose distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ satisfies $E\left[\kappa(\mathbf{R}_i)^2\right], E\left[\kappa(\mathbf{S}_i)^2\right] \geq \sigma_{\mathsf{Con}}^2, \forall i \in [P]$, there exists a distribution $\mathbb{P}_{\mathbf{A},\mathbf{B}}$ satisfying Assumption 2.1 such that*

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) \geq \frac{\eta^2}{\left(1 + \frac{\eta}{\sigma_{\mathsf{Con}}^2}\right)^2}.$$

The above converse is translated to a bound in terms of ϵ-the DP parameter. It is worth noting that the converse does not necessarily assume Laplace distributions, and is applicable to *any* distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ that satisfies ϵ-DP.

Theorem 3.4.

$$\mathsf{LMSE}^*(P, 2, \epsilon, 1) \geq \frac{\eta^2}{\left(\eta \frac{e^\epsilon - 1}{L^2} + 1\right)^2} \tag{3.3}$$

Proof Sketch (Missing details in Appendix A). We observe that lower bounding σ_{Con}^2 in Theorem 3.3 should give us the desired relation. We show a proof sketch for the scalar case $L = 1$, the general proof is given in Appendix A . Without loss of generality assume k^{th} node attains the minimum variance $\sigma_{\mathsf{Con}}^2 = \mathbb{E}[R_k^2]$. Denote the pdf of R_k to be p_{R_k}. Then we can write from the DP Definition 2.1,

$$\frac{p_{R_k}(r-1)}{p_{R_k}(r)} \leq e^\epsilon.$$

$$\sigma_{\mathsf{Con}}^2 = \int_{-\infty}^{\infty} r^2 p_{R_k}(r) dr \tag{3.4}$$

$$\geq e^{-\epsilon} \int_{-\infty}^{\infty} r^2 p_{R_k}(r-1) dr \tag{3.5}$$

$$= e^{-\epsilon} \int_{-\infty}^{\infty} (r+1)^2 p_{R_k}(r) dr \tag{3.6}$$

$$= e^{-\epsilon}(\sigma_{\mathsf{Con}}^2 + 1) \tag{3.7}$$

$$\implies \sigma_{\mathsf{Con}}^2 \geq \frac{1}{e^\epsilon - 1} \tag{3.8}$$

Substituting the above in the result of Theorem 3.3 gives us the desired relation. □

(a) LMSE vs ϵ

Figure 1: The privacy-accuracy trade-offs of Theorems 3.2 and 3.4. The LMSE is plotted in logarithmic scale.

Remark 2. For the case where we are multiplying scalars ($L = 1$), the smallest and largest singular values are the same. For this case, Theorems 3.1 and 3.3 combine to give a tight characterization of the trade-off between LMSE and the standard deviation of additive the noise, that is:

$$\inf_{\Gamma} \mathsf{LMSE}(\Gamma) = \frac{\eta^2}{\left(1 + \frac{\eta}{\sigma^2}\right)^2}. \tag{3.9}$$

where the infimum is over all coding schemes Γ whose probability distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ satisfies[7]

$$\mathbb{E}\left[R_i^2\right], \mathbb{E}\left[S_i^2\right] \geq \sigma^2, \forall i \in [P].$$

IV. CONCLUSION

This work opens a new direction via the search of codes that optimize privacy-utility trade-off for secure multiparty computing. There are several open questions motivated by our work. First, the study of optimal code design for $t \geq 1$ is a natural open question. While an achievable scheme can be developed along the same lines as our paper by assessing the Shamir secret sharing scheme with arbitrary close evaluation points, development of a tight characterization (even for the scalar case of $L = 1$, with the standard deviation measure on the privacy loss) is an open problem. Second, our coding schemes do not assume shared randomness, that is, they assume $\mathbf{R}, \mathbf{S}$ are statistically independent. The question of whether shared randomness can improve the accuracy-privacy trade-off is an interesting open question. Finally, because our schemes require evaluation points that are arbitrarily close to each other, the computation nodes need to perform computations at a high level of precision. This can involve hidden computation and storage costs (see, a similar phenomenon in [27], [28]). An explicit characterization of these hidden costs is an interesting area of future work.

---

[7] We have dropped the boldface notation in the subsequent equation to indicate that the quantities are scalars

References

[1] D. Evans, V. Kolesnikov, and M. Rosulek, "A pragmatic introduction to secure multi-party computation," *Foundations and Trends® in Privacy and Security*, vol. 2, no. 2-3, 2017.

[2] M. Ben-Or, S. Goldwasser, and A. Wigderson, "Completeness theorems for non-cryptographic fault-tolerant distributed computation," in *Proceedings of the Twentieth Annual ACM Symposium on Theory of Computing*, ser. STOC '88. New York, NY, USA: Association for Computing Machinery, 1988, p. 1–10. [Online]. Available: https://doi.org/10.1145/62212.62213

[3] A. Shamir, "How to share a secret," *Commun. ACM*, vol. 22, no. 11, p. 612–613, nov 1979. [Online]. Available: https://doi.org/10.1145/359168.359176

[4] Q. Yu, S. Li, N. Raviv, S. M. M. Kalan, M. Soltanolkotabi, and S. A. Avestimehr, "Lagrange coded computing: Optimal design for resiliency, security, and privacy," in *Proceedings of the Twenty-Second International Conference on Artificial Intelligence and Statistics*, ser. Proceedings of Machine Learning Research, K. Chaudhuri and M. Sugiyama, Eds., vol. 89. PMLR, 16–18 Apr 2019, pp. 1215–1225. [Online]. Available: https://proceedings.mlr.press/v89/yu19b.html

[5] H. B. McMahan, G. Andrew, U. Erlingsson, S. Chien, I. Mironov, N. Papernot, and P. Kairouz, "A general approach to adding differential privacy to iterative training procedures," *arXiv preprint arXiv:1812.06210*, 2018.

[6] "Tensorflow privacy." [Online]. Available: https://www.tensorflow.org/responsible_ai/privacy/guide

[7] C. Dwork and A. Roth, "The algorithmic foundations of differential privacy," *Foundations and Trends® in Theoretical Computer Science*, vol. 9, no. 3–4, pp. 211–407, 2014. [Online]. Available: http://dx.doi.org/10.1561/0400000042

[8] C. Dwork, N. Kohli, and D. Mulligan, "Differential privacy in practice: Expose your epsilons!" *Journal of Privacy and Confidentiality*, vol. 9, no. 2, 2019.

[9] S. Yuan, M. Shen, I. Mironov, and A. C. A. Nascimento, "Practical, label private deep learning training based on secure multiparty computation and differential privacy," Cryptology ePrint Archive, Report 2021/835, 2021. [Online]. Available: https://ia.cr/2021/835

[10] M. Pettai and P. Laud, "Combining differential privacy and secure multiparty computation," in *Proceedings of the 31st Annual Computer Security Applications Conference*, ser. ACSAC 2015. New York, NY, USA: Association for Computing Machinery, 2015, p. 421–430. [Online]. Available: https://doi.org/10.1145/2818000.2818027

[11] X. He, A. Machanavajjhala, C. Flynn, and D. Srivastava, "Composing differential privacy and secure computation: A case study on scaling private record linkage," in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '17. New York, NY, USA: Association for Computing Machinery, 2017, p. 1389–1406. [Online]. Available: https://doi.org/10.1145/3133956.3134030

[12] J. Böhler and F. Kerschbaum, "Secure multi-party computation of differentially private median," in *29th USENIX Security Symposium (USENIX Security 20)*. USENIX Association, Aug. 2020, pp. 2147–2164. [Online]. Available: https://www.usenix.org/conference/usenixsecurity20/presentation/boehler

[13] M. Soleymani, H. Mahdavifar, and A. S. Avestimehr, "Analog lagrange coded computing," *IEEE Journal on Selected Areas in Information Theory*, vol. 2, no. 1, pp. 283–295, 2021. [Online]. Available: https://doi.org/10.1109/JSAIT.2021.3056377

[14] H. Akbari-Nodehi and M. A. Maddah-Ali, "Secure coded multi-party computation for massive matrix operations," *IEEE Transactions on Information Theory*, vol. 67, no. 4, pp. 2379–2398, 2021.

[15] W.-T. Chang and R. Tandon, "On the capacity of secure distributed matrix multiplication," in *2018 IEEE Global Communications Conference (GLOBECOM)*, 2018, pp. 1–6.

[16] R. G. L. D'Oliveira, S. El Rouayheb, and D. Karpuk, "Gasp codes for secure distributed matrix multiplication," *IEEE Transactions on Information Theory*, vol. 66, no. 7, pp. 4038–4050, 2020.

[17] Z. Jia and S. A. Jafar, "On the capacity of secure distributed batch matrix multiplication," *IEEE Transactions on Information Theory*, vol. 67, no. 11, pp. 7420–7437, 2021.

[18] Z. Chen, Z. Jia, Z. Wang, and S. A. Jafar, "Gcsa codes with noise alignment for secure coded multi-party batch matrix multiplication," *IEEE Journal on Selected Areas in Information Theory*, vol. 2, no. 1, pp. 306–316, 2021.

[19] M. Fahim and V. R. Cadambe, "Numerically Stable Polynomially Coded Computing," *IEEE Transactions on Information Theory*, p. 1, 2021. [Online]. Available: https://ieeexplore.ieee.org/document/9319171

[20] N. Agarwal and K. Singh, "The price of differential privacy for online learning," in *International Conference on Machine Learning*. PMLR, 2017, pp. 32–40.

[21] J. C. Duchi, M. I. Jordan, and M. J. Wainwright, "Local privacy and statistical minimax rates," in *2013 IEEE 54th Annual Symposium on Foundations of Computer Science*. IEEE, 2013, pp. 429–438.

[22] S. Asoodeh, M. Aliakbarpour, and F. P. Calmon, "Local differential privacy is equivalent to contraction of e_γ-divergence," *arXiv preprint arXiv:2102.01258*, 2021.

[23] Q. Geng and P. Viswanath, "The optimal noise-adding mechanism in differential privacy," *IEEE Transactions on Information Theory*, vol. 62, no. 2, pp. 925–951, 2016. [Online]. Available: https://doi.org/10.1109/TIT.2015.2504967

[24] Q. Geng, W. Ding, R. Guo, and S. Kumar, "Tight analysis of privacy and utility tradeoff in approximate differential privacy," in *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*, ser. Proceedings of Machine Learning Research, S. Chiappa and R. Calandra, Eds., vol. 108. PMLR, 26–28 Aug 2020, pp. 89–99. [Online]. Available: https://proceedings.mlr.press/v108/geng20a.html

[25] P. Kairouz, S. Oh, and P. Viswanath, "Secure multi-party differential privacy," in *Advances in Neural Information Processing Systems*, C. Cortes, N. Lawrence, D. Lee, M. Sugiyama, and R. Garnett, Eds., vol. 28. Curran Associates, Inc., 2015. [Online]. Available: https://proceedings.neurips.cc/paper/2015/file/a01610228fe998f515a72dd730294d87-Paper.pdf

[26] ——, "Differentially private multi-party computation," in *2016 Annual Conference on Information Science and Systems (CISS)*, 2016, pp. 128–132. [Online]. Available: https://doi.org/10.1109/CISS.2016.7460489

[27] H. Jeong, A. Devulapalli, V. R. Cadambe, and F. P. Calmon, "ϵ-approximate coded matrix multiplication is nearly twice as efficient as exact multiplication," *IEEE Journal on Selected Areas in Information Theory*, vol. 2, no. 3, pp. 845–854, 2021. [Online]. Available: https://doi.org/10.1109/JSAIT.2021.3099811

[28] J. Wang, Z. Jia, and S. A. Jafar, "Price of precision in coded distributed matrix multiplication: A dimensional analysis," in *2021 IEEE Information Theory Workshop (ITW)*. IEEE, 2021, pp. 1–6.

In this section we provide proofs for the achievability and converse theorems stated in Section III.

A. Achievability

We prove the Theorems 3.1 and 3.2 in the following.

Let the subset $\mathcal{S} = \{i, j\}$. As stated in Theorem 3.1, we show an achievable scheme for a subset of distributions $(\mathbf{R}, \mathbf{S})$, where node i gets evaluations as,

$$\tilde{\mathbf{A}}_i = \mathbf{A} + \mathbf{\Lambda}u_i, \tilde{\mathbf{B}}_i = \mathbf{B} + \mathbf{\Theta}u_i.$$

We pick $\mathbb{P}_{\mathbf{\Lambda},\mathbf{\Theta}}$ such that $\mathbb{E}[\mathbf{\Lambda}] = \mathbb{E}[\mathbf{\Theta}] = 0$ and $|u_i| \geq \sigma_{\mathsf{Ach}}$. For convenience of illustration we drop the dependence on $\mathcal{S}$ for decoding weights, i.e., $w_{1,\mathcal{S}}, w_{2,\mathcal{S}}$ will be written as w_1, w_2 respectively. Thus,

$$\widehat{\mathbf{C}}_{\mathcal{S}} = w_1 \tilde{\mathbf{A}}_i \tilde{\mathbf{B}}_i + w_2 \tilde{\mathbf{A}}_j \tilde{\mathbf{B}}_j.$$

Let

$$\bar{E}(w_1, w_2, u_i, u_j) \triangleq (w_1 + w_2 - 1)^2\eta^2 + 2\eta(w_1 u_i + w_2 u_j)^2 + (w_1 u_i^2 + w_2 u_j^2)^2.$$

We will now prove an upper bound on $\bar{E}(w_1, w_2, u_i, u_j)$ which will be used to prove Theorem 3.1.

Lemma A.1. *For any $\Delta > 0$ there exist some $w_1^*, w_2^* \in \mathbb{R}$ and $u_i, u_j \in \mathbb{R}$ satisfying $|u_i|, |u_j| \geq \sigma_{\mathsf{Ach}}$ such that*

$$\bar{E}(w_1^*, w_2^*, u_i, u_j) \leq \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2})^2} + \Delta.$$

Proof. We find w_1^*, w_2^* by minimizing $\bar{E}(w_1, w_2, u_i, u_j)$ over w_1, w_2. Equating the Jacobian of $\bar{E}(w_1, w_2, u_i, u_j)$ with respect to w_1, w_2 to 0 and solving for w_1, w_2 gives,

$$w_1^* = \frac{\eta(-u_i u_j^2 - u_j^3 - 2\eta u_j)}{(u_i - u_j)\left(2u_i^2 u_j^2 + \eta(u_i + u_j)^2 + 2\eta^2\right)} \quad \text{(A.1)}$$

$$w_2^* = \frac{\eta(u_i^2 u_j + u_i^3 + 2\eta u_i)}{(u_i - u_j)\left(2u_i^2 u_j^2 + \eta(u_i + u_j)^2 + 2\eta^2\right)}. \quad \text{(A.2)}$$

Observe that $\bar{E}(w_1, w_2, u_i, u_j)$ is a convex quadratic in w_1, w_2. Substituting w_1^*, w_2^* in $\bar{E}(w_1, w_2, u_i, u_j)$.

$$\min_{w_1, w_2 \in \mathbb{R}} \bar{E}(w_1, w_2, u_i, u_j) = \bar{E}(w_1^*, w_2^*, u_i, u_j)$$

$$= \frac{2\eta^2 u_i^2 u_j^2}{2u_i^2 u_j^2 + \eta(u_i + u_j)^2 + 2\eta^2} \quad \text{(A.3)}$$

Now, let $\epsilon_1 = \frac{1}{\sigma_{\mathsf{Ach}}^2} - \frac{\eta}{u_i^2}$, $\epsilon_2 = \frac{1}{\sigma_{\mathsf{Ach}}^2} - \frac{\eta}{u_j^2}$ and $\epsilon_3 = (u_i - u_j)^2$.

To bring about the desired relation, we do the following,

$$\frac{1}{\eta^2}\left(1 + \frac{\eta}{u_i^2}\right)\left(1 + \frac{\eta}{u_j^2}\right) - \frac{1}{\bar{E}(w_1^*, w_2^*, u_i, u_j)}$$

$$= \frac{1}{2}\frac{(u_i - u_j)^2}{\eta u_i^2 u_j^2} \quad \text{(A.4)}$$

$$\leq \frac{1}{2\eta\sigma_{\mathsf{Ach}}^4}\epsilon_3$$

$$\frac{1}{\eta^2}\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2} - \epsilon_1\right)\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2} - \epsilon_2\right) - \frac{1}{\bar{E}(w_1^*, w_2^*, u_i, u_j)}$$

$$\leq \frac{1}{2\eta\sigma_{\mathsf{Ach}}^4}\epsilon_3 \quad \text{(A.5)}$$

$$\frac{1}{\eta^2}\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2}\right)^2 - \frac{1}{\bar{E}(w_1^*, w_2^*, u_i, u_j)}$$

$$\leq \frac{1}{\eta^2}\left(\left(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2}\right)(\epsilon_1 + \epsilon_2) - \epsilon_1\epsilon_2 + \frac{\eta}{2\sigma_{\mathsf{Ach}}^4}\epsilon_3\right) \triangleq \frac{h}{\eta^2}. \quad \text{(A.6)}$$

$$\implies \bar{E}(w_1^*, w_2^*, u_i, u_j) \leq \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2})^2 - h}. \quad \text{(A.7)}$$

Remark 3. We ideally want h to be close to 0. Observe that picking $|u_i|$ and $|u_j|$ close to $\sqrt{\eta}\sigma_{\mathsf{Ach}}$ makes $\epsilon_1, \epsilon_2, \epsilon_3$ close to 0 which gives h close to 0.

Taylor series expansion about $h = 0$ gives

$$\bar{E}(w_1^*, w_2^*, u_i, u_j) \leq \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2})^2} + O(h). \quad \text{(A.8)}$$

Taking $\Delta = O(h)$ gives,

$$\implies \bar{E}(w_1^*, w_2^*, u_i, u_j) \leq \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2})^2} + \Delta. \quad \text{(A.9)}$$

$\square$

We now use the above result to prove Theorem 3.1

Proof of Theorem 3.1.

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) = \mathbb{E}[||\mathbf{AB} - \widehat{\mathbf{C}}_{\mathcal{S}}||_F^2]$$

$$= \mathbb{E}[||w_1(\mathbf{AB} + (\mathbf{\Lambda B} + \mathbf{A\Theta})u_i + \mathbf{\Lambda\Theta}u_i^2)$$
$$+ w_2(\mathbf{AB} + (\mathbf{\Lambda B} + \mathbf{A\Theta})u_j + \mathbf{\Lambda\Theta}u_j^2) - \mathbf{AB}||_F^2] \quad \text{(A.10)}$$

$$= \mathbb{E}[||(w_1 + w_2 - 1)\mathbf{AB} + (w_1 u_i + w_2 u_j)(\mathbf{\Lambda B} + \mathbf{A\Theta})$$
$$+ (w_1 u_i^2 + w_2 u_j^2)\mathbf{\Lambda\Theta}||_F^2] \quad \text{(A.11)}$$

Since $\mathbf{A}, \mathbf{B}, \mathbf{\Lambda}, \mathbf{\Theta}$ are independent and $\mathbf{\Lambda}, \mathbf{\Theta}$ are zero mean, the cross products vanish,

$$= (w_1 + w_2 - 1)^2 \mathbb{E}[||\mathbf{AB}||_F^2]$$
$$+ (w_1 u_i + w_2 u_j)^2 (\mathbb{E}[||\mathbf{\Lambda B}||_F^2]$$
$$+ \mathbb{E}[||\mathbf{A\Theta}||_F^2]) + (w_1 u_i^2 + w_2 u_j^2)^2 \mathbb{E}[||\mathbf{\Lambda\Theta}||_F^2] \quad \text{(A.12)}$$

From system model and from theorem statement we know $\mathbb{E}||\mathbf{A}||_F^2, \mathbb{E}[||\mathbf{B}||_F^2] \leq \eta, \mathbb{E}[||\mathbf{\Lambda}||_F^2] = \mathbb{E}[||\mathbf{\Theta}||_F^2] = 1$ and using sub-multiplicative property of Frobenius norm and independence of $\mathbf{A}, \mathbf{B}, \mathbf{\Lambda}, \mathbf{\Theta}$ gives,

$$\leq (w_1 + w_2 - 1)^2 \eta^2 + 2\eta(w_1 u_i + w_2 u_j)^2 + (w_1 u_i^2 + w_2 u_j^2)^2 \quad \text{(A.13)}$$

From Corollary A.1

$$\leq \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\mathsf{Ach}}^2})^2} + \Delta \quad \text{(A.14)}$$

Since the analysis was applied to arbitrary set $\{i, j\}$, the analysis holds for every $N = 2$ element subset $\mathcal{S}$. $\quad\square$

Making the distributional assumptions given in Theorem 3.2, we provide a relation between σ_{Ach} and ϵ.

Proof of Theorem 3.2. Observe that for the i^{th} node if $(\mathbf{\Lambda}_i)_{m,n} \sim \text{Laplace}(0, \frac{1}{\sqrt{2}L})$, then the distribution of $(\mathbf{R}_i)_{m,n}$ is,

$$f_{(\mathbf{R}_i)_{m,n}}(z) = \frac{L}{\sqrt{2}|u_i|} \exp\left(-\sqrt{2}L \frac{|z|}{|u_i|}\right) \quad \text{(A.15)}$$

Similarly,

$$f_{(\mathbf{S}_i)_{m,n}}(z) = \frac{L}{\sqrt{2}|u_i|} \exp\left(-\sqrt{2}L \frac{|z|}{|u_i|}\right) \quad \text{(A.16)}$$

Let's evaluate the ratio in Definition 2.1 at a point $\bar{\mathbf{Y}} = \begin{bmatrix} \bar{\mathbf{Y}}_0 \in \mathbb{R}^{L \times L} \\ \bar{\mathbf{Y}}_1 \in \mathbb{R}^{L \times L} \end{bmatrix}$ i.e., let $\mathcal{A} = \{\bar{\mathbf{Y}}\}$. Let $\mathbf{X}_l = \begin{bmatrix} \mathbf{A}_l \\ \mathbf{B}_l \end{bmatrix}$.

$$\frac{\mathbb{P}(\mathbf{Y}^{(0)} \in \mathcal{A})}{\mathbb{P}(\mathbf{Y}^{(1)} \in \mathcal{A})} = \frac{\mathbb{P}_{\mathbf{R}_i}(\bar{\mathbf{Y}} - \mathbf{X}_0)}{\mathbb{P}_{\mathbf{R}_i}(\bar{\mathbf{Y}} - \mathbf{X}_1)} \quad \text{(A.17)}$$

$$= \prod_{k,l \in [L]} \frac{\exp\left(-\sqrt{2}L \frac{|(\bar{\mathbf{Y}}_0)_{m,n} - (\mathbf{A}_0)_{m,n}|}{|u_i|}\right)}{\exp\left(-\sqrt{2}L \frac{|(\bar{\mathbf{Y}}_0)_{m,n} - (\mathbf{A}_1)_{m,n}|}{|u_i|}\right)}$$

$$\frac{\exp\left(-\sqrt{2}L \frac{|(\bar{\mathbf{Y}}_1)_{m,n} - (\mathbf{B}_0)_{m,n}|}{|u_i|}\right)}{\exp\left(-\sqrt{2}L \frac{|(\bar{\mathbf{Y}}_1)_{m,n} - (\mathbf{B}_1)_{m,n}|}{|u_i|}\right)} \quad \text{(A.18)}$$

$$\leq \prod_{k,l \in [L]} \exp\left(\sqrt{2}L \frac{|(\mathbf{A}_0)_{m,n} - (\mathbf{A}_1)_{m,n}|}{|u_i|}\right)$$

$$\exp\left(\sqrt{2}L \frac{|(\mathbf{B}_0)_{m,n} - (\mathbf{B}_1)_{m,n}|}{|u_i|}\right) \quad \text{(A.19)}$$

Since $|u_i| \geq \sigma_{\mathsf{Ach}}$ and letting $\beta_{m,n} = |(\mathbf{A}_0)_{m,n} - (\mathbf{A}_1)_{m,n}| + |(\mathbf{B}_0)_{m,n} - (\mathbf{B}_1)_{m,n}|$,

$$\leq \prod_{k,l \in [L]} \exp\left(\frac{\sqrt{2}L}{\sigma_{\mathsf{Ach}}} \beta_{m,n}\right) \quad \text{(A.20)}$$

$$= \exp\left(\frac{\sqrt{2}L}{\sigma_{\mathsf{Ach}}} \sum_{k,l \in [L]} \beta_{m,n}\right) \quad \text{(A.21)}$$

$$\leq \exp\left(\frac{2\sqrt{2}L^3}{\sigma_{\mathsf{Ach}}} ||\mathbf{X}_0 - \mathbf{X}_1||_{\max}\right) \quad \text{(A.22)}$$

$$\leq \exp\left(\frac{2\sqrt{2}L^3}{\sigma_{\mathsf{Ach}}}\right) = \exp(\epsilon) \quad \text{(A.23)}$$

Thus, we take $\sigma_{\mathsf{Ach}}^2 = \frac{8L^6}{\epsilon^2}$. We obtain the last equation by letting $\sigma_{\mathsf{Ach}} = \frac{2\sqrt{2}L^3}{\epsilon}$. Finally, plugging this σ_{Ach} in equation (3.1) completes the proof. $\quad\square$

B. Converse

In this subsection we prove the Theorems 3.3 and 3.4.

Proof of Theorem 3.3.
Consider a (P, N) coding scheme Γ which satisfies $\mathbb{E}(\kappa(\mathbf{R}_i)^2) \geq \sigma_{\mathsf{Con}}^2$ and $\mathbb{E}(\kappa(\mathbf{S}_i)^2) \geq \sigma_{\mathsf{Con}}^2$ for all i. Similar to achievability section we do the analysis for some two nodes i and j, which can be applied to any two nodes. Thus let the subset $\mathcal{S} = \{i, j\}$.

Pick a distribution $\mathbb{P}_{\mathbf{A}, \mathbf{B}}$ such that

$$\mathbf{A} = \begin{bmatrix} a & 0 & \dots & 0 \\ \vdots & & & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix}$$

$$\mathbf{B} = \begin{bmatrix} b & 0 & \dots & 0 \\ \vdots & & & \vdots \\ 0 & 0 & \dots & 0 \end{bmatrix}$$

where $\mathbb{E}[a] = \mathbb{E}[b] = 0, \mathbb{E}[a^2] = \mathbb{E}[b^2] = \eta$. Observe that $\mathbb{E}[||\mathbf{A}||_F^2], \mathbb{E}[||\mathbf{B}||_F^2] = \eta$ and $\mathbb{E}[\mathbf{A}] = \mathbb{E}[\mathbf{B}] = 0$.

Again for convenience of illustration we drop the dependence on $\mathcal{S}$ for decoding weights, i.e., $w_{1,\mathcal{S}}, w_{2,\mathcal{S}}$ will be written as w_1, w_2 respectively. Thus,

$$\widehat{\mathbf{C}}_{\mathcal{S}} = w_1 \tilde{\mathbf{A}}_i \tilde{\mathbf{B}}_i + w_2 \tilde{\mathbf{A}}_j \tilde{\mathbf{B}}_j.$$

In the following we use MATLAB inspired notation in several places. For example, for a matrix $\mathbf{A}$, the notation $\mathbf{A}[i, :]$ would mean a row vector constituting i^{th} row of $\mathbf{A}$, while $\mathbf{A}[:, i]$ is a column vector constituting i^{th} column of $\mathbf{A}$. The notation $\mathbf{A}[i, j]$ means $(i, j)^{\text{th}}$ element of $\mathbf{A}$. The notation $\mathbf{A}[i, j : k]$ means a row vector constituting elements j through k including k of i^{th} row. Also for any two matrices $\mathbf{A}$ and $\mathbf{B}$ horizontal contatenation operation is denoted by the notation $[\mathbf{A} \ \mathbf{B}]$.

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) = \mathbb{E}[||\mathbf{AB} - \widehat{\mathbf{C}}_{\mathcal{S}}||_F^2]$$
$$= \mathbb{E}[||(w_1 + w_2 - 1)\mathbf{AB} + (w_1\mathbf{R}_i + w_2\mathbf{R}_j)\mathbf{B}$$
$$+ \mathbf{A}(w_1\mathbf{S}_i + w_2\mathbf{S}_j) + (w_1\mathbf{R}_i\mathbf{S}_i + w_2\mathbf{R}_j\mathbf{S}_j)||_F^2] \quad (A.24)$$

Since $\mathbf{A}, \mathbf{B}, [\mathbf{R}_i\ \mathbf{R}_j], [\mathbf{S}_i\ \mathbf{S}_j]$ are independent and $\mathbf{A}, \mathbf{B}$ are zero mean, the cross products vanish,

$$= (w_1 + w_2 - 1)^2 \mathbb{E}[||\mathbf{AB}||_F^2] + \mathbb{E}[||(w_1\mathbf{R}_i + w_2\mathbf{R}_j)\mathbf{B}||_F^2]$$
$$+ \mathbb{E}[||\mathbf{A}(w_1\mathbf{S}_i + w_2\mathbf{S}_j)||_F^2]$$
$$+ \mathbb{E}[||w_1\mathbf{R}_i\mathbf{S}_i + w_2\mathbf{R}_j\mathbf{S}_j||_F^2] \quad (A.25)$$

First, let us look at the last term of (A.25),

$$\mathbb{E}[||w_1\mathbf{R}_i\mathbf{S}_i + w_2\mathbf{R}_j\mathbf{S}_j||_F^2] \quad (A.26)$$
$$\geq \mathbb{E}[(w_1\mathbf{R}_i[1,:]\mathbf{S}_i[:,1] + w_2\mathbf{R}_j[1,:]\mathbf{S}_j[:,1])^2] \quad (A.27)$$
$$= \mathsf{Var}(w_1\mathbf{R}_i[1,:]\mathbf{S}_i[:,1] + w_2\mathbf{R}_j[1,:]\mathbf{S}_j[:,1])$$
$$+ \mathbb{E}[(w_1\mathbf{R}_i[1,:]\mathbf{S}_i[:,1] + w_2\mathbf{R}_j[1,:]\mathbf{S}_j[:,1])]^2 \quad (A.28)$$

Clearly taking $\mathbb{E}[\mathbf{R}_i[1,:]] = \mathbb{E}[\mathbf{S}_j[:,1]] = 0$, reduces LMSE. Let the autocorrelation matrix of $\mathbf{R}_i[1,:]$ be $\mathbf{K}_{\mathbf{R}_i[1,:]}$. We now derive a bound on $\mathbf{K}_{\mathbf{R}_i[1,:]}$. For any unit vector $\mathbf{x} \in \mathbb{R}^{L \times 1}$,

$$\sigma_{\mathsf{Con}}^2 \leq \mathbb{E}[\kappa(\mathbf{R}_i)^2] \quad (A.29)$$
$$\overset{(1)}{\leq} \mathbb{E}[([1\ 0\ \ldots\ 0]\mathbf{R}_i\mathbf{x})^2] \quad (A.30)$$
$$= \mathbb{E}[\mathbf{x}^\top \mathbf{R}_i[1,:]^\top \mathbf{R}_i[1,:]\mathbf{x}] \quad (A.31)$$
$$= \mathbf{x}^\top \mathbf{K}_{\mathbf{R}_i[1,:]}\mathbf{x}. \quad (A.32)$$
$$\implies \mathbf{K}_{\mathbf{R}_i[1,:]} \succcurlyeq \sigma_{\mathsf{Con}}^2 \mathbf{I} \quad (A.33)$$

Inequality (1) comes from the fact that for any matrix $\mathbf{\Lambda}$, $\kappa(\mathbf{\Lambda}) \leq \mathbf{x}^\top \mathbf{\Lambda}\mathbf{y}$ for all unit vectors $\mathbf{x}, \mathbf{y}$. Without loss of generality we take $\mathbf{R}_i, \mathbf{R}_j, \mathbf{S}_i, \mathbf{S}_j$ to be gaussian random matrices, that is, elements are non-iid gaussian distributed. We are able to do choose a distribution without losing generality because LMSE depends only on second order statistics, this is seen by expanding the Frobenius norms in (A.25) and observing that it consists of only variances and covariances.

Let $u_i \in \mathbb{R}$. Now choose a random variable $\bar{r}_i$ such that

$$u_i\bar{r}_i = \mathbf{R}_i[1,1] - \mathbb{E}[\mathbf{R}_i[1,1] \,|\, \mathbf{R}_i[1,2], \ldots, \mathbf{R}_i[1,L]].$$

Essentially $u_i\bar{r}_i$ is the MMSE estimation error of $\mathbf{R}_i[1,1]$ when $\mathbf{R}_i[1,2], \ldots, \mathbf{R}_i[1,L]$ are observed. Since $\mathbf{R}_i[1,:]$ is a gaussian vector, this is a linear estimator, therefore we can write, for some $\bar{\mathbf{w}} \in \mathbb{R}^{L-1 \times 1}$ and $c \in \mathbb{R}$,

$$\mathbb{E}[\mathbf{R}_i[1,1] \,|\, \mathbf{R}_i[1,2], \ldots, \mathbf{R}_i[1,L]] = \mathbf{R}_i[1, 2:L]\bar{\mathbf{w}} + c$$

$$u_i\bar{r}_i = [\mathbf{R}_i[1,:]\ 1] \begin{bmatrix} 1 \\ -\bar{\mathbf{w}} \\ -c \end{bmatrix}.$$

Let $\tilde{\mathbf{w}} = \begin{bmatrix} 1 \\ -\bar{\mathbf{w}} \end{bmatrix}$.

$$u_i^2\bar{r}_i^2 = [\tilde{\mathbf{w}}^\top\ -c] \begin{bmatrix} \mathbf{R}_i[1,:]^\top \mathbf{R}_i[1,:] & \mathbf{R}_i[1,:]^\top \\ \mathbf{R}_i[1,:] & 1 \end{bmatrix} \begin{bmatrix} \tilde{\mathbf{w}} \\ -c \end{bmatrix} \quad (A.34)$$
$$= \tilde{\mathbf{w}}^\top \mathbf{R}_i[1,:]^\top \mathbf{R}_i[1,:]\tilde{\mathbf{w}} - 2\mathbf{R}_i[1,:]\tilde{\mathbf{w}}c + c^2 \quad (A.35)$$

Since $\mathbb{E}[\mathbf{R}_i[1,:]] = 0$, $c = 0$ and therefore,

$$u_i^2\mathbb{E}[\bar{r}_i^2] = \mathbb{E}[(\mathbf{R}_i[1,:]\bar{\mathbf{w}})^2] \geq \sigma_{\mathsf{Con}}^2 \quad (A.36)$$

Observe that $\mathbb{E}[\bar{r}_i] = 0$ and thus without loss of generality we can take $\mathbb{E}[\bar{r}_i^2] = 1$, then $u_i^2 \geq \sigma_{\mathsf{Con}}^2$. Let

$$\mathbf{z}_1 = \mathbf{R}_i[1,:]^\top - u_i[\bar{r}_i\ 0\ \ldots\ 0]^\top.$$

Then from orthogonality principle of MMSE estimation theory we see that $\mathbb{E}[\bar{r}_i\mathbf{z}_1] = \mathbf{0}_{L \times 1}$. Using a similar process for $\mathbf{S}_j[:,1]$ let

$$\mathbf{z}_2 = \mathbf{S}_j[:,1] - v_j[\bar{s}_j\ 0\ \ldots\ 0]^\top,$$

where $v_j^2 \geq \sigma_{\mathsf{Con}}^2$, $\mathbb{E}[\bar{s}_j^2] = 1$, and from orthogonality principle we see that $\mathbb{E}[\bar{s}_j\mathbf{z}_2] = \mathbf{0}_{L \times 1}$.

Let us look at each of four terms in (A.25).

$$\mathbb{E}[||\mathbf{AB}||_F^2] = \mathbb{E}[(ab)^2] = \eta^2 \quad (A.37)$$

$$\mathbb{E}[||(w_1\mathbf{R}_i + w_2\mathbf{R}_j)\mathbf{B}||_F^2] \geq \mathbb{E}[||(w_1\mathbf{R}_i + w_2\mathbf{R}_j)\mathbf{B}||_2^2]$$
$$\overset{(1)}{\geq} \mathbb{E}[([1\ 0\ \ldots\ 0](w_1\mathbf{R}_i + w_2\mathbf{R}_j)\mathbf{B}[1\ 0\ \ldots\ 0]^\top)^2] \quad (A.38)$$
$$\geq \mathbb{E}[((w_1\mathbf{R}_i[1,:] + w_2\mathbf{R}_j[1,:])\mathbf{B}[:,1])^2] \quad (A.39)$$
$$= \mathbb{E}[((w_1\mathbf{R}_i[1,1] + w_2\mathbf{R}_j[1,1])b)^2] \quad (A.40)$$
$$= \mathbb{E}[(w_1\mathbf{R}_i[1,1] + w_2\mathbf{R}_j[1,1])^2]\mathbb{E}[b^2] \quad (A.41)$$
$$\overset{(2)}{\geq} \mathbb{E}[\bar{r}_i(w_1\mathbf{R}_i[1,1] + w_2\mathbf{R}_j[1,1])]^2\ \eta \quad (A.42)$$
$$= \eta(w_1u_i + w_2u_j)^2 \quad (A.43)$$

where $u_j = \mathbb{E}[\bar{r}_i\mathbf{R}_j[1,1]]$. Inequality (1) comes from the fact that for any matrix $\mathbf{\Lambda}, ||\mathbf{\Lambda}||_2 \geq \mathbf{x}^\top \mathbf{\Lambda}\mathbf{y}$ for all unit vectors $\mathbf{x}, \mathbf{y}$. Inequality (2) is from Cauchy-Schwarz inequality.

Similarly,

$$\mathbb{E}[||\mathbf{A}(w_1\mathbf{S}_i + w_2\mathbf{S}_j)||_F^2] \geq \eta(w_1v_i + w_2v_j)^2,$$

where $v_i = \mathbb{E}[\bar{s}_j\mathbf{S}_i[1,1]]$.

For the last term in (A.25).

$$\mathbb{E}[||w_1\mathbf{R}_i\mathbf{S}_i + w_2\mathbf{R}_j\mathbf{S}_j||_F^2] \quad (A.44)$$
$$\geq \mathbb{E}[(w_1\mathbf{R}_i[1,:]\mathbf{S}_i[:,1] + w_2\mathbf{R}_j[1,:]\mathbf{S}_j[:,1])^2] \quad (A.45)$$
$$\geq \mathbb{E}[\bar{r}_i(w_1\mathbf{R}_i[1,:]\mathbf{S}_i[:,1] + w_2\mathbf{R}_j[1,:]\mathbf{S}_j[:,1])\bar{s}_j]^2 \quad (A.46)$$
$$= (w_1u_i\mathbb{E}[\mathbf{S}_i[1,1]\bar{s}_j] + w_2\mathbb{E}[\bar{r}_i\mathbf{R}_j[1,1]]v_j)^2 \quad (A.47)$$
$$= (w_1u_iv_i + u_jv_j)^2. \quad (A.48)$$

Therefore, combining all four terms (A.25) is lower bounded

by,

$$(w_1 + w_2 - 1)^2\eta^2 + \eta(w_1 u_i + w_2 u_j)^2$$
$$+ \eta(w_1 v_i + w_2 v_j)^2 + (w_1 u_i v_i + w_2 u_j v_j)^2 \quad \text{(A.49)}$$
$$\triangleq ||\mathbf{M}||_F^2 \quad \text{(A.50)}$$

where $\mathbf{M}$ is defined to be:

$$\mathbf{M} = \left(w_1 \begin{bmatrix} \sqrt{\eta} \\ u_i \end{bmatrix} \begin{bmatrix} \sqrt{\eta} & v_i \end{bmatrix} + w_2 \begin{bmatrix} \sqrt{\eta} \\ u_j \end{bmatrix} \begin{bmatrix} \sqrt{\eta} & v_j \end{bmatrix} - \begin{bmatrix} \eta & 0 \\ 0 & 0 \end{bmatrix} \right).$$
$$\text{(A.51)}$$

Multiplying by $\begin{bmatrix} -v_j \\ \sqrt{\eta} \end{bmatrix}$ on both sides and taking ℓ_2-norm,

$$\left\| \mathbf{M} \begin{bmatrix} -v_j \\ \sqrt{\eta} \end{bmatrix} \right\|_2^2 = (w_1(v_i - v_j) + v_j)^2\eta^2 + w_1^2(v_i - v_j)^2 u_i^2 \eta$$
$$\text{(A.52)}$$

The above is a convex quadratic equation in w_1 and is minimized at,

$$w_1^* = \frac{\eta v_j}{(\eta + u_i^2)(v_i - v_j)}.$$

Substituting w_1^* in equation (A.52) and simplifying we write,

$$\frac{\eta^2 v_j^2 u_i^2}{\eta + u_i^2} \le \left\| \mathbf{M} \begin{bmatrix} -v_j \\ \sqrt{\eta} \end{bmatrix} \right\|_2^2 \quad \text{(A.53)}$$

$$\le ||\mathbf{M}||_2^2 (\eta + v_j^2) \quad \text{(A.54)}$$

$$\implies ||\mathbf{M}||_F^2 \ge ||\mathbf{M}||_2^2 \ge \eta^2 \frac{u_i^2}{\eta + u_i^2} \frac{v_j^2}{\eta + v_j^2} \quad \text{(A.55)}$$

$$\ge \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\text{Con}}^2})^2} \quad \text{(A.56)}$$

Therefore,

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) \ge \frac{\eta^2}{(1 + \frac{\eta}{\sigma_{\text{Con}}^2})^2} \quad \text{(A.57)}$$

$\square$

Similar as in achievability, we now prove a lower bound on σ_{Con}^2 and use it the above result.

Proof of Theorem 3.4. Consider a (P, N) coding scheme Γ which satisfies $\mathbb{E}(\kappa(\mathbf{R}_i)^2) \ge \sigma_{\text{Con}}^2$ and $\mathbb{E}(\kappa(\mathbf{S}_i)^2) \ge \sigma_{\text{Con}}^2$. We aim to lower bound σ_{Con}^2 for a given ϵ, i.e., we say that given some ϵ we cannot use σ_{Con}^2 less than some quantity and still satisfy ϵ-DP. Thus, we assume worst case inputs and derive σ_{Con}^2 achievable for that worst case scenario. We further use this to show that LMSE below some quantity is not achievable given an ϵ. Without loss of generality assume that $\sigma_{\text{Con}}^2 = \mathbb{E}[\kappa(\mathbf{R}_k)^2]$, it suffices to concentrate on k^{th} node. Let $\mathbf{X}_l = \begin{bmatrix} \mathbf{A}_l \\ \mathbf{B}_l \end{bmatrix}$. And let,

$$\mathbf{Y}^{(0)} = \mathbf{X}_0 + \begin{bmatrix} \mathbf{R}_k \\ \mathbf{S}_k \end{bmatrix},$$

$$\mathbf{Y}^{(1)} = \mathbf{X}_1 + \begin{bmatrix} \mathbf{R}_k \\ \mathbf{S}_k \end{bmatrix},$$

From differential privacy Definition 2.1, for any subset $\mathcal{A} \subset \mathbb{R}^{2L \times L}$ for any $||\mathbf{X}_0 - \mathbf{X}_1||_{\max} \le 1$ it is true that

$$\frac{P(\mathbf{Y}^{(1)} \in \mathcal{A})}{P(\mathbf{Y}^{(0)} \in \mathcal{A})} \le e^\epsilon \text{ and } \frac{P(\mathbf{Y}^{(0)} \in \mathcal{A})}{P(\mathbf{Y}^{(1)} \in \mathcal{A})} \le e^\epsilon.$$

The worst case inputs have $||\mathbf{X}_0 - \mathbf{X}_1||_{\max} = 1$. Without loss of generality we pick $\mathbf{X}_0 = \mathbf{0}_{2L \times L}$ and $\mathbf{X}_1 = \mathbf{1}_{2L \times L}$, where $\mathbf{0}_{2L \times L}$ and $\mathbf{1}_{2L \times L}$ are $2L \times L$ matrices with all elements 0's and 1's respectively. For some $\bar{\mathbf{R}} \in \mathbb{R}^{L \times L}$, pick $\mathcal{A} = \left\{ \begin{bmatrix} \bar{\mathbf{R}} \\ \mathbb{R}^{L \times L} \end{bmatrix} \right\}$. Denote the pdf of $\mathbf{R}_k$ to be $p_{\mathbf{R}_k}$ and $\mathbf{1}$ an $L \times L$ matrix with all elements to be 1. Then by independence of $\mathbf{R}_k$ and $\mathbf{S}_k$ and from our choice of $\mathcal{A}$, we can effectively ignore the role of $\mathbf{S}_k$ and rewrite the above as,

$$\frac{p_{\mathbf{R}_k}(\bar{\mathbf{R}} - \mathbf{1})}{p_{\mathbf{R}_k}(\bar{\mathbf{R}})} \le e^\epsilon \text{ and } \frac{p_{\mathbf{R}_k}(\bar{\mathbf{R}})}{p_{\mathbf{R}_k}(\bar{\mathbf{R}} - \mathbf{1})} \le e^\epsilon.$$

Using post processing property [7] of differential privacy we can use $\kappa(\cdot)$ map without violating ϵ-DP. Denote the pdf of $\kappa(\mathbf{R}_k)$ to be $p_{\kappa(\mathbf{R}_k)}$. After $\kappa(\cdot)$ mapping we have two cases,

1) $\kappa(\bar{\mathbf{R}} - \mathbf{1}) \le \kappa(\bar{\mathbf{R}})$: Take

$$\frac{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}} - \mathbf{1}))}{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}}))} \le e^\epsilon.$$

$$\implies P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}} - \mathbf{1})) \le e^\epsilon P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}}))$$

Using Weyl's inequality for singular values $\kappa(\bar{\mathbf{R}} - \mathbf{1}) \ge \kappa(\bar{\mathbf{R}}) - ||\mathbf{1}||_2$, we write,

$$P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}}) - ||\mathbf{1}||_2) \le e^\epsilon P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}}))$$

$$\implies \frac{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}}) - ||\mathbf{1}||_2)}{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}}))} \le e^\epsilon.$$

Let $\kappa(\bar{\mathbf{R}}) = r$ and we know that $||\mathbf{1}||_2 = L$, we write,

$$\frac{p_{\kappa(\mathbf{R}_k)}(r - L)}{p_{\kappa(\mathbf{R}_k)}(r)} \le e^\epsilon.$$

2) $\kappa(\bar{\mathbf{R}} - \mathbf{1}) \ge \kappa(\bar{\mathbf{R}})$: Take

$$\frac{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}}))}{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}} - \mathbf{1}))} \le e^\epsilon.$$

$$\implies P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}})) \le e^\epsilon P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}} - \mathbf{1}))$$

Using Weyl's inequality for singular values , $\kappa(\bar{\mathbf{R}} - \mathbf{1}) \le \kappa(\bar{\mathbf{R}}) + ||\mathbf{1}||_2$, we write,

$$P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}} - \mathbf{1}) - ||\mathbf{1}||_2) \le e^\epsilon P(\kappa(\mathbf{R}_k) \le \kappa(\bar{\mathbf{R}} - \mathbf{1}))$$

$$\implies \frac{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}} - \mathbf{1}) - ||\mathbf{1}||_2)}{p_{\kappa(\mathbf{R}_k)}(\kappa(\bar{\mathbf{R}} - \mathbf{1}))} \le e^\epsilon.$$

Let $\kappa(\bar{\mathbf{R}} - \mathbf{1}) = r$ and we know that $||\mathbf{1}||_2 = L$, we write,

$$\frac{p_{\kappa(\mathbf{R}_k)}(r - L)}{p_{\kappa(\mathbf{R}_k)}(r)} \le e^\epsilon.$$

Thus, we have shown that the above relation is true for any

r, in turn true for arbitrary choice of $\bar{\mathbf{R}}$. Now,

$$\sigma_{\mathsf{Con}}^2 = \mathbb{E}[\kappa(\mathbf{R}_k)^2] = \int_{-\infty}^{\infty} r^2 p_{\kappa(\mathbf{R}_k)}(r)dr \tag{A.58}$$

$$\geq e^{-\epsilon} \int_{-\infty}^{\infty} r^2 p_{\kappa(\mathbf{R}_k)}(r-L)dr \tag{A.59}$$

$$= e^{-\epsilon} \int_{-\infty}^{\infty} (r+L)^2 p_{\kappa(\mathbf{R}_k)}(r)dr \tag{A.60}$$

$$= e^{-\epsilon}(\sigma_{\mathsf{Con}}^2 + L^2 + 2L\mathbb{E}[\kappa(\mathbf{R}_k)]) \tag{A.61}$$

$$\geq e^{-\epsilon}(\sigma_{\mathsf{Con}}^2 + L^2) \tag{A.62}$$

$$\implies \sigma_{\mathsf{Con}}^2 \geq \frac{L^2}{e^\epsilon - 1} \tag{A.63}$$

Substituting the above in (A.57) gives us the desired relation.
$\square$

APPENDIX B
ACCURACY PRIVACY TRADE-OFF FOR $t = 1, N = 1$ CASE.

Theorem B.1. *Let $\mathbf{\Lambda}, \mathbf{\Theta}$ be $L \times L$ independent zero-mean random matrices with i.i.d. entries each with a variance of $1/L^2$. For any $\Delta, \sigma_{\mathsf{Ach}} > 0$ there exist scalars $u_i, i \in [P]$ with $|u_i| \geq \sigma_{\mathsf{Ach}}$ such that, if $\mathbf{R}_i = u_i\mathbf{\Lambda}, \mathbf{S}_i = u_i\mathbf{\Theta}, i \in [P]$, then there is a $(P, N = 1)$ coding scheme Γ with distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ such that, for every $\mathbb{P}_{\mathbf{A},\mathbf{B}}$ satisfying Assumption 2.1,*

$$\mathsf{LMSE}(\Gamma) \leq 2\eta\sigma_{\mathsf{Ach}}^2 + \sigma_{\mathsf{Ach}}^4 + \Delta. \tag{B.1}$$

Theorem B.2. *For any $(P, N = 1)$ code Γ whose distribution $\mathbb{P}_{\mathbf{R},\mathbf{S}}$ satisfies $E\left[\kappa(\mathbf{R}_i)^2\right], E\left[\kappa(\mathbf{S}_i)^2\right] \geq \sigma_{\mathsf{Con}}^2, \forall i \in [P]$, there exists a distribution $\mathbb{P}_{\mathbf{A},\mathbf{B}}$ satisfying Assumption 2.1 such that*

$$\mathsf{LMSE}_{\mathcal{S}}(\Gamma) \geq 2\eta\sigma_{\mathsf{Con}}^2 + \sigma_{\mathsf{Con}}^4.$$

We will not provide rigourous proofs of these theorems, since this case is not the focus of our paper, but observe that this case is obtained by taking $w_1 = 1$ and $w_2 = 0$ in Equations (A.13) and (A.49). To obtain the relation to ϵ, we use Theorems 3.2 and 3.4.